

ICS 212: Operating Systems

Lecture 1

Instructor:

Dr. Nilotpai Chakraborty

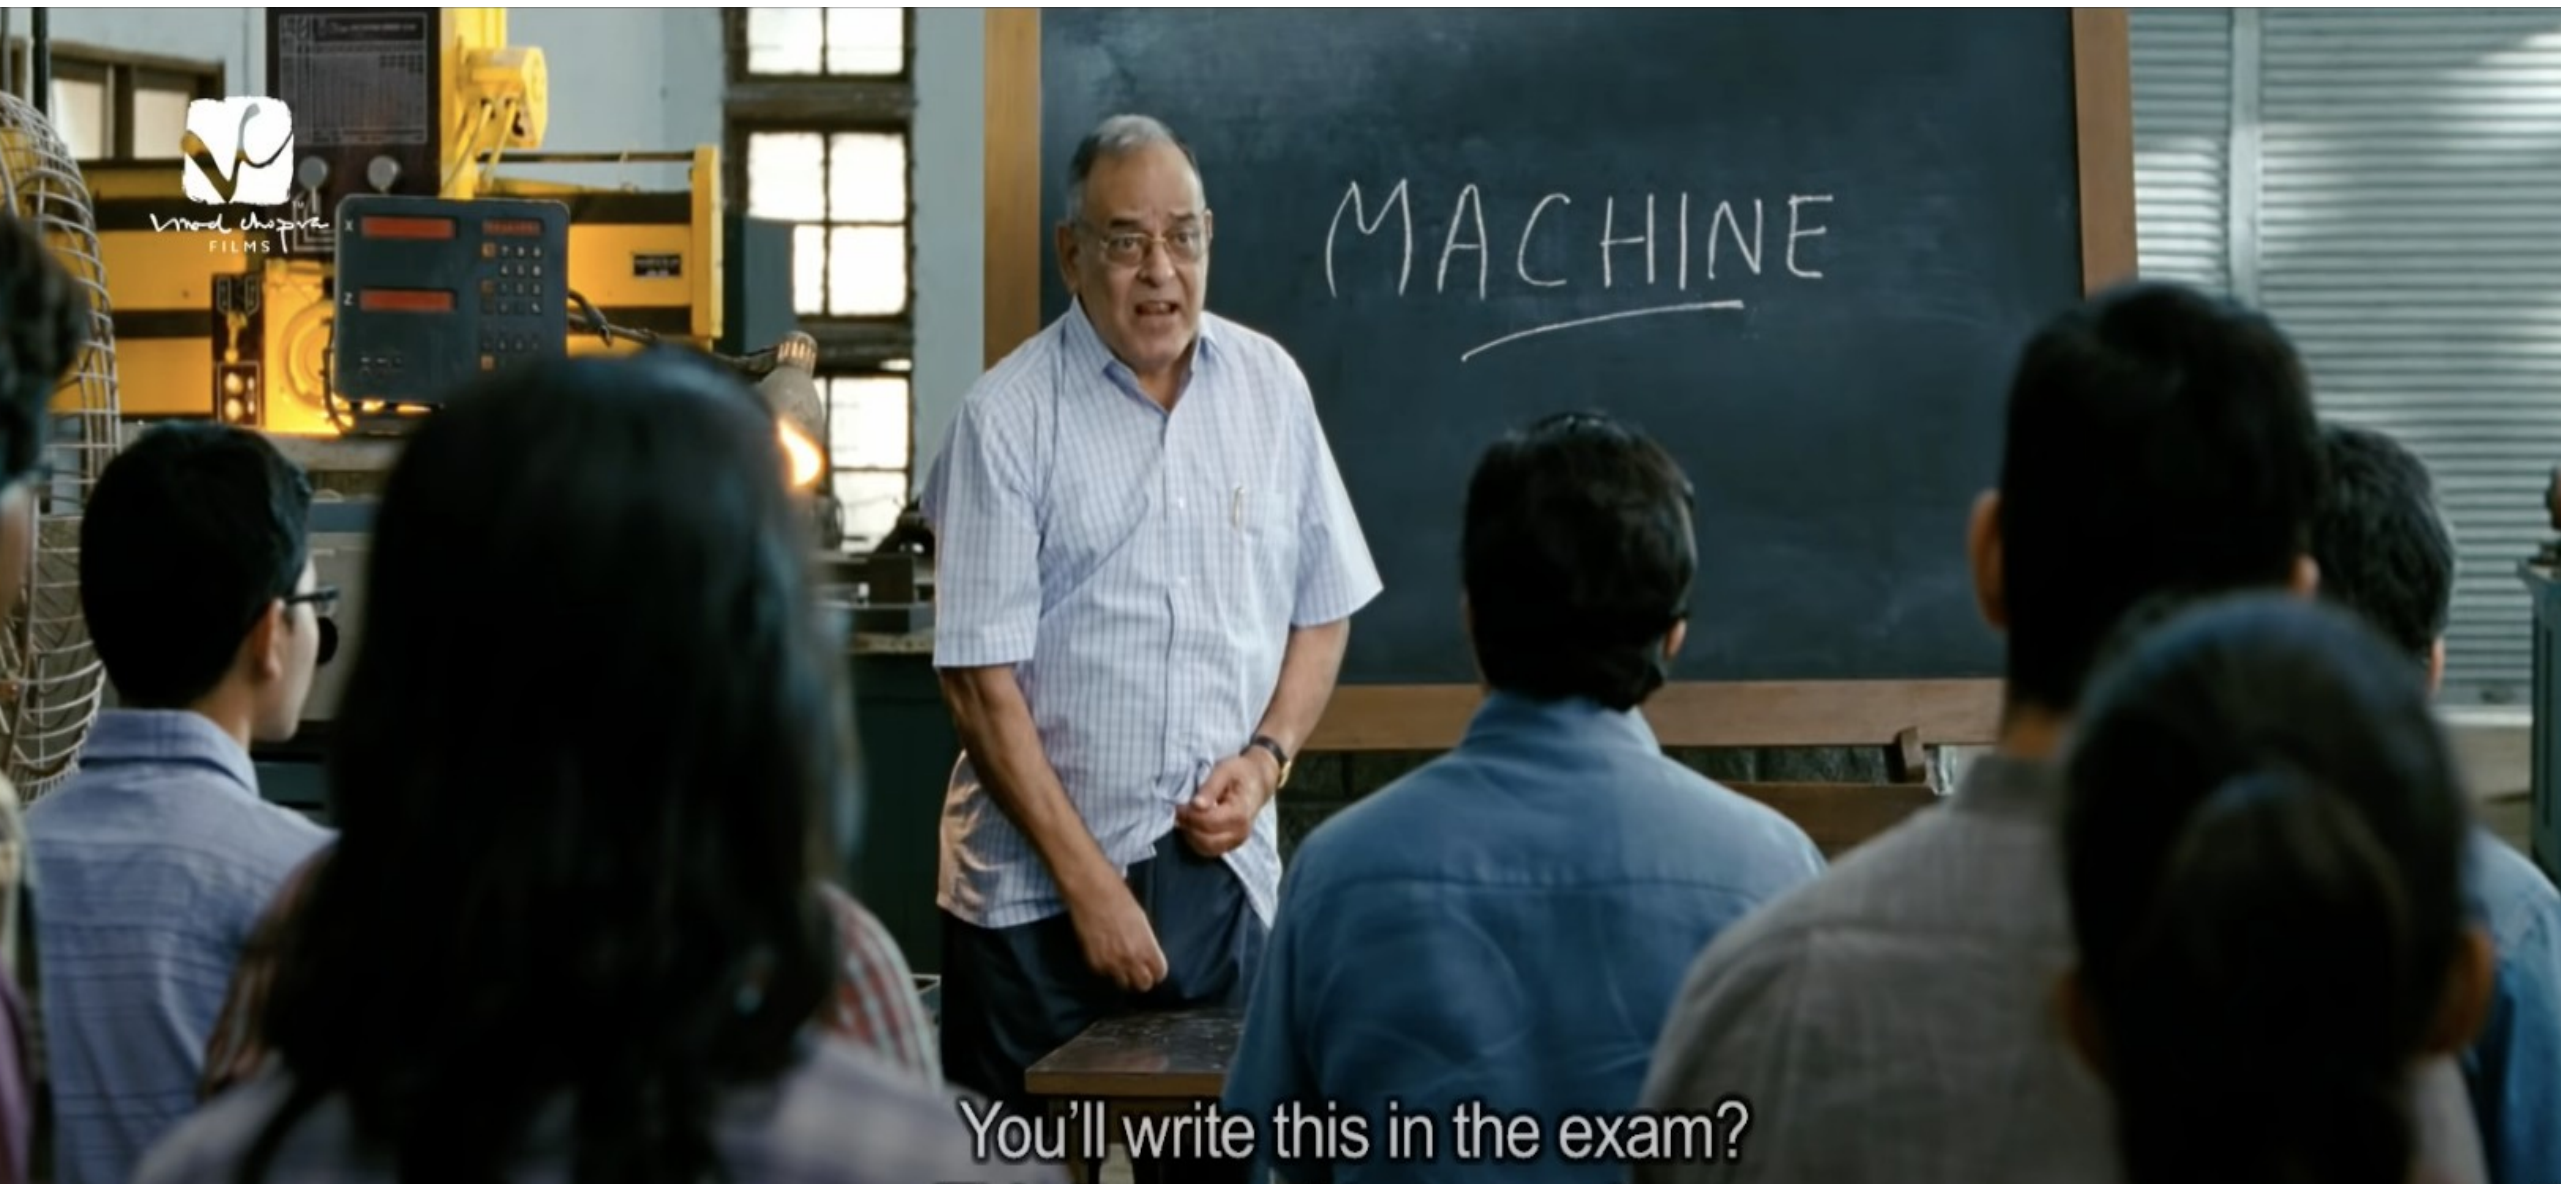
Indian Institute of Information Technology Kottayam

Books and reading materials

- **William Stallings**, *Operating systems: Internals & design principles*, Pearson, Seventh edition, 2014.
- **Silberschatz, Galvin, Gagne**, *Operating System Concepts*, Wiley, Ninth Edition, 2016.
- Other online resources and papers.

Let's define OS..





You'll write this in the exam?

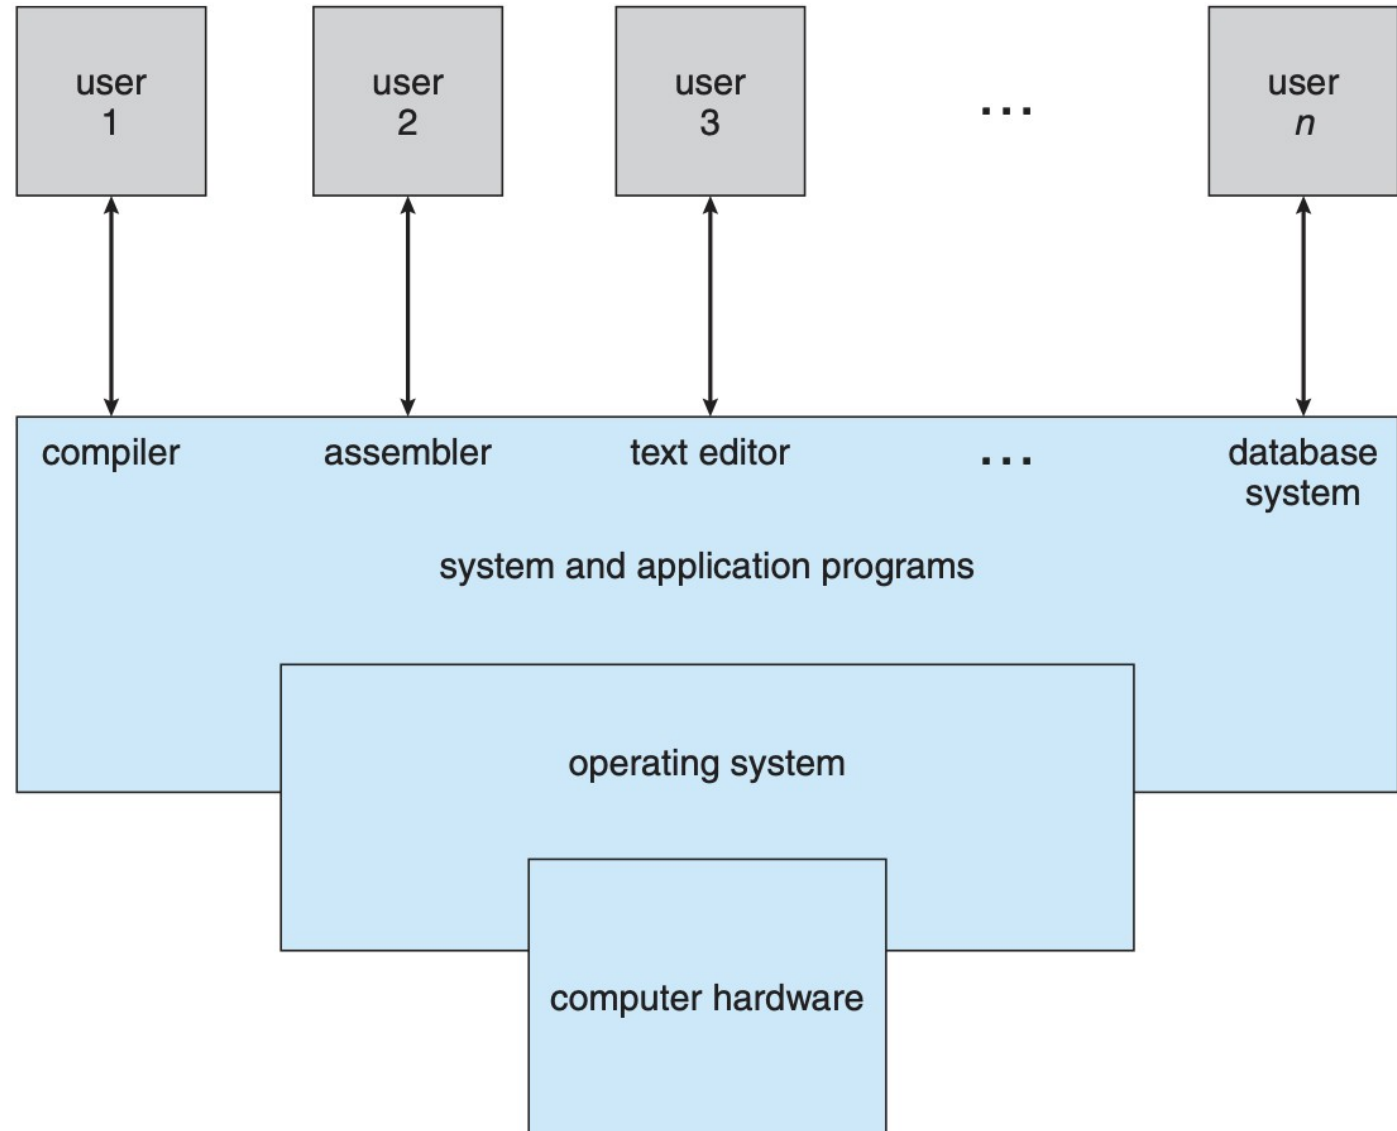
From the treatise

- **Stallings**: “An **Operating System** is a program that controls the execution of application programs and acts as an interface between the applications and the computer hardware.”
- **Galvin**: “An **operating system** is a program that manages a computer’s hardware. It also provides a basis for application programs and acts as an intermediary between the computer user and the computer hardware.”
- **Objectives**: *convenience*, *efficiency*, *evolving*

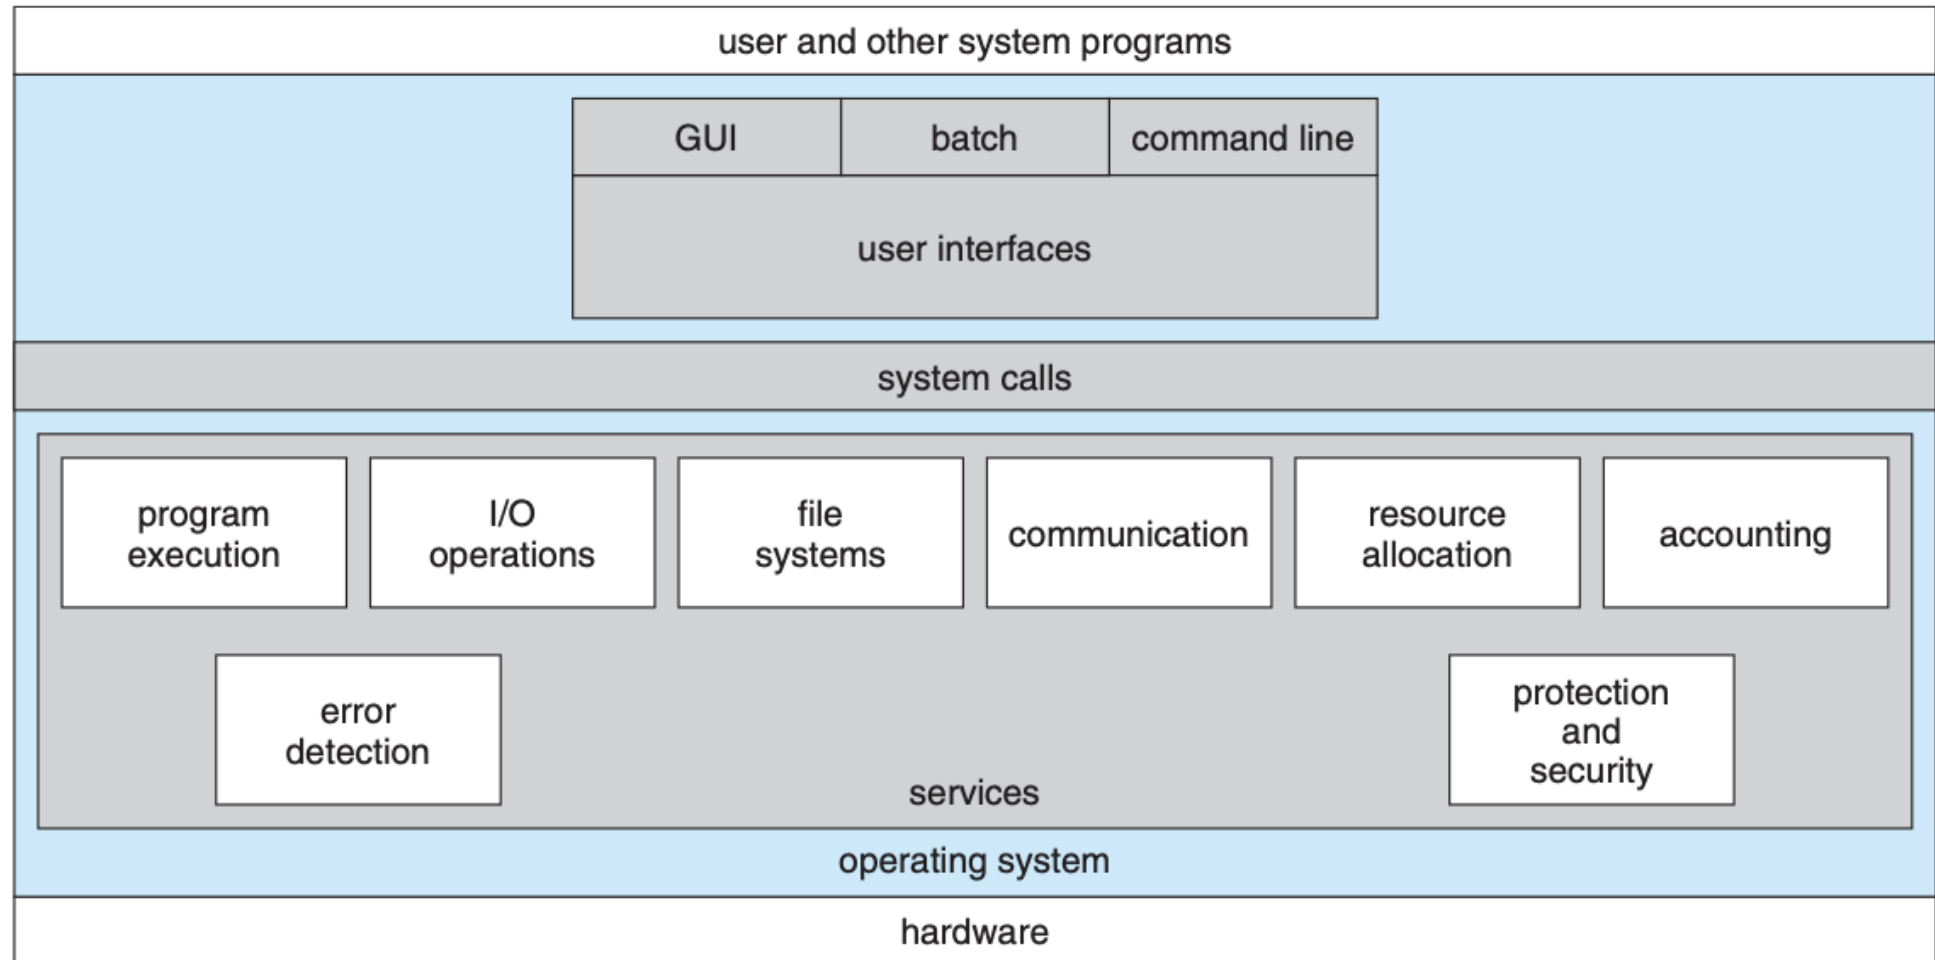
OS as per functionality

- Provides *interface* between hardware and software
- A *resource manager* that allows programs/users to share the hardware resources
- A *controller* that controls execution of programs to prevent errors and improper use of the computer
- *Creates a virtual machine* that is easier to program than the raw hardware
- Overall services by OS: *program development and execution, device control and access, error control, accounting, API, etc.*

Abstract View



OS services



Over the years..

- **Single user computers:**
 - One user at a time using the console
 - One function at a time for the computer
- **Batch processing:**
 - Executes multiple jobs in batches
 - Jobs are submitted on cards or tapes, which are in turn batched sequentially
- **Multiprogrammed Batch Systems:**
 - Overlap I/O jobs and computing jobs
 - CPU can execute another job while waiting for another job
- **Multiprogramming:**
 - Several programs can run simultaneously
 - OS manages interactions

Thank You!

ICS 212: Operating Systems


Lecture 2

Instructor:

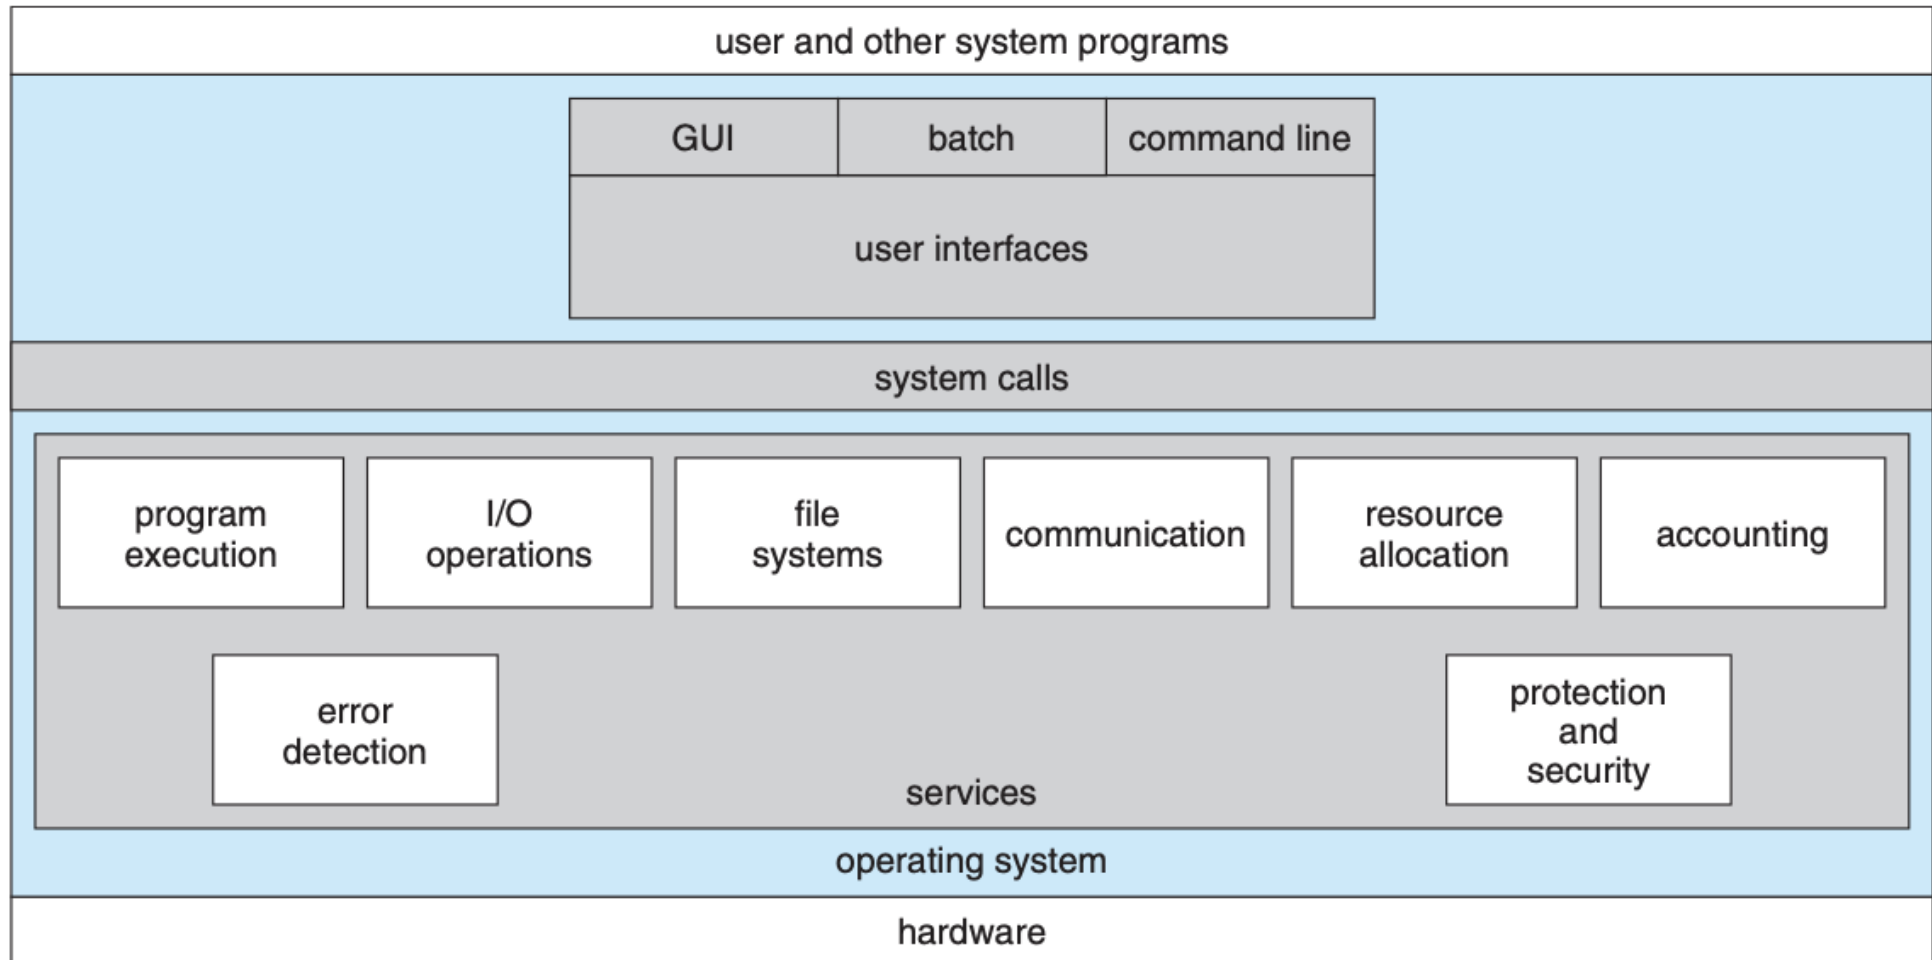
Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

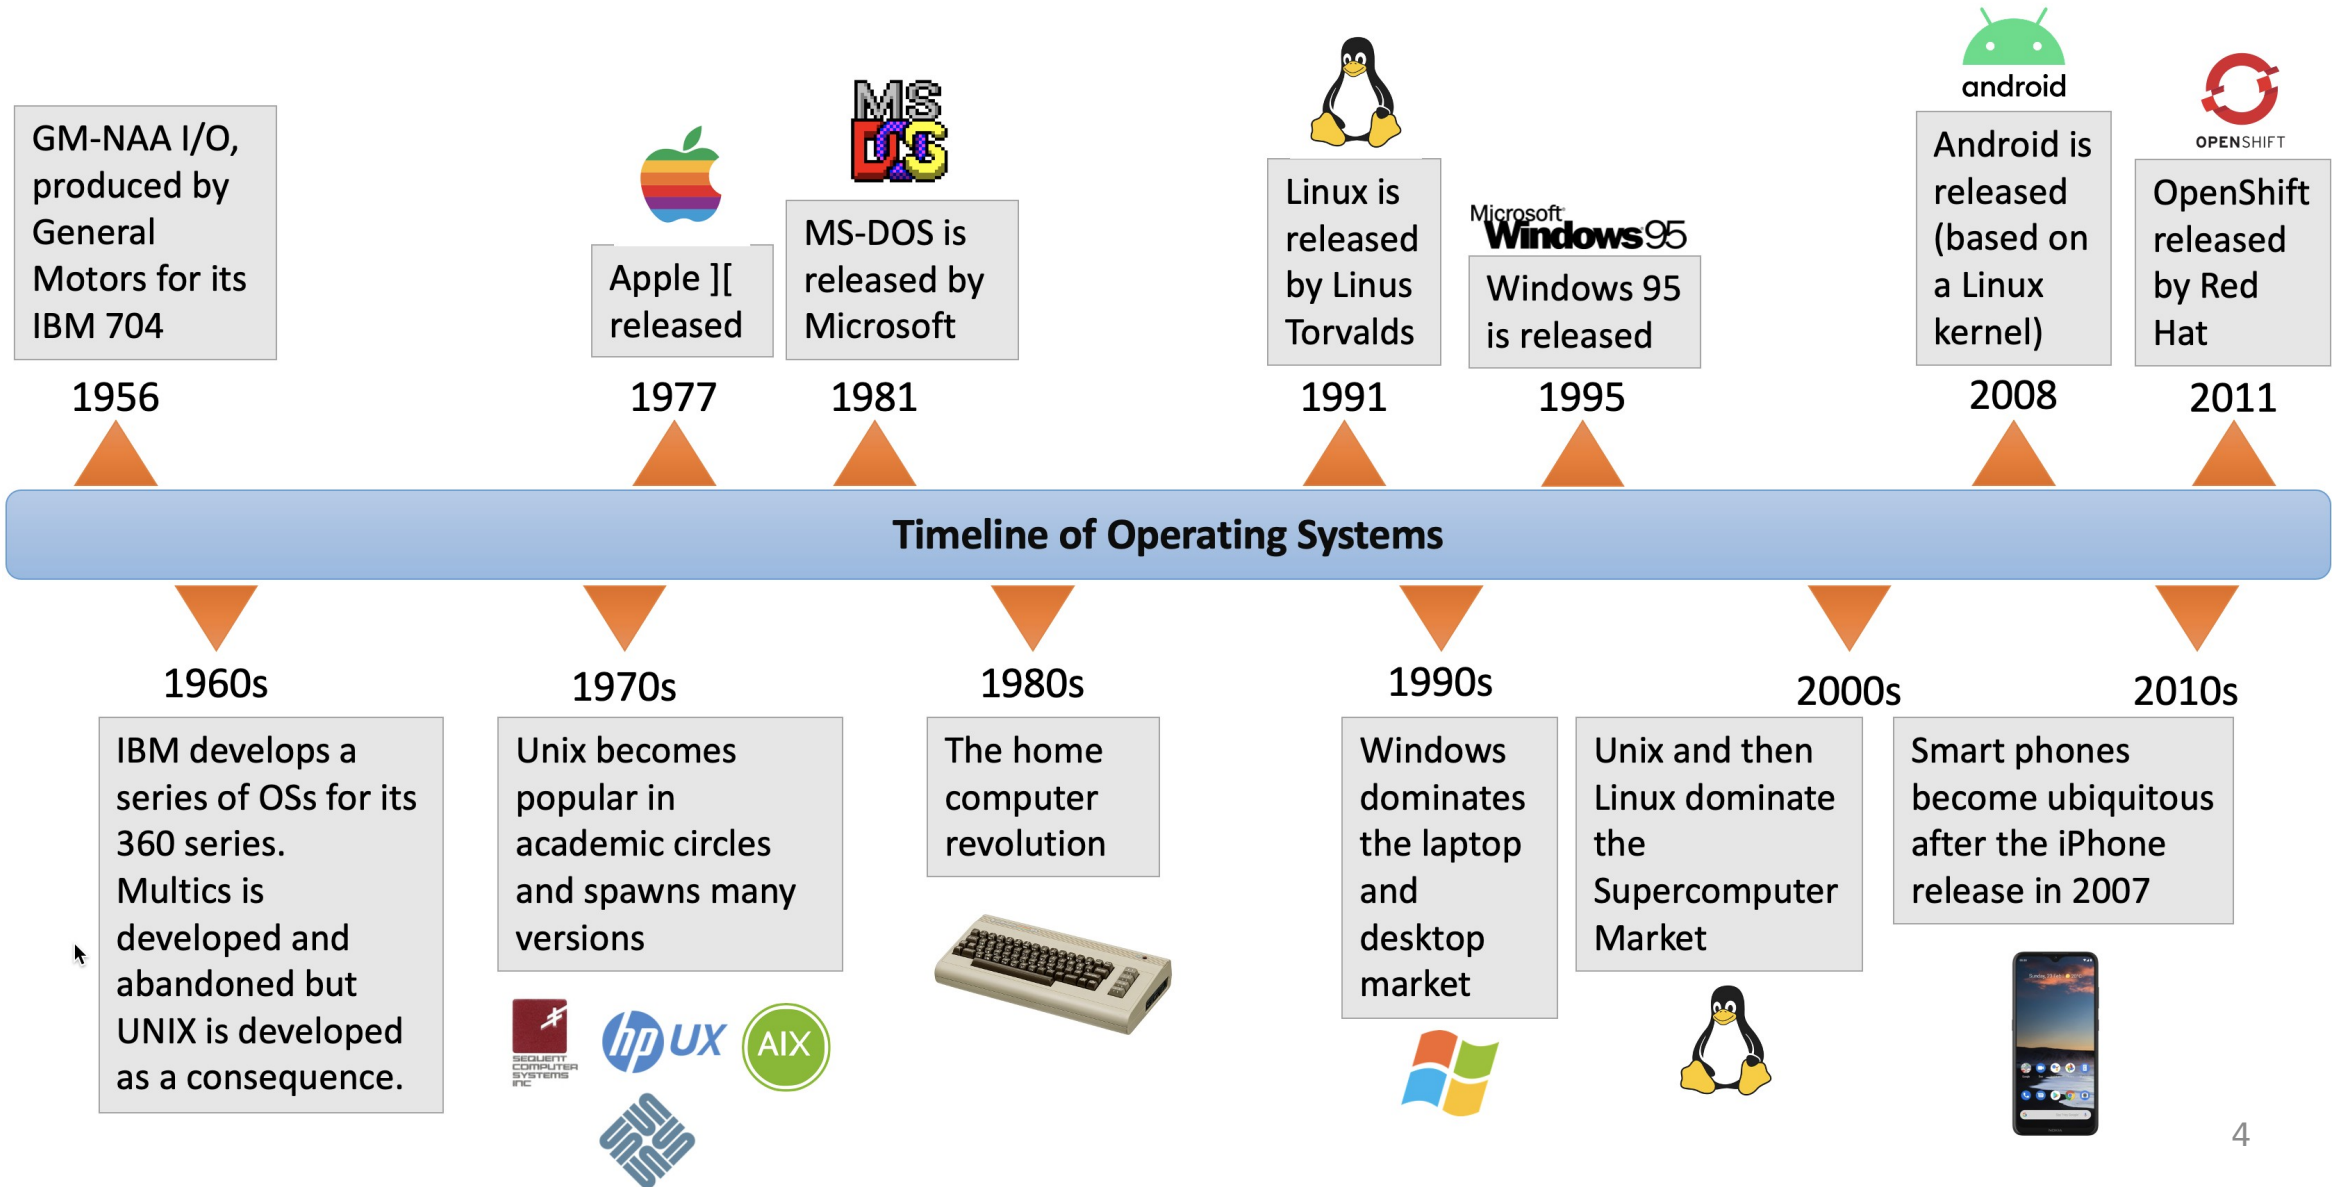
Abstract View



OS services



OS Timeline



Computer System Organization

Architectural features

OS service	Hardware support
Interrupts	Interrupt vectors
Protection	Kernel/user mode, protected instructions, base/limit registers
System calls	Trap instructions, Trap vectors
I/O	Interrupts and memory mapping
Virtual memory	Translation lookaside buffers

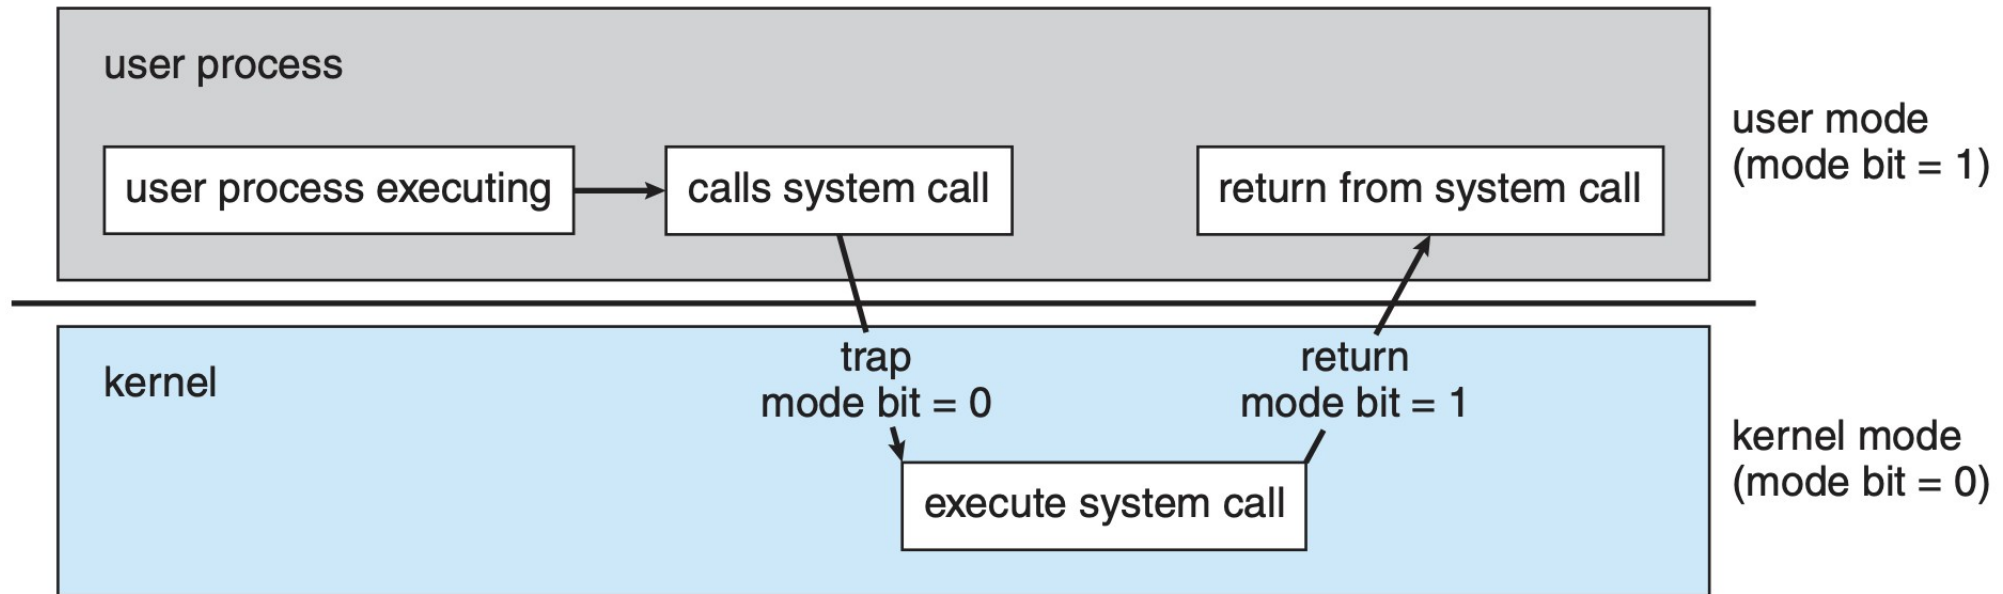
Functionalities of Interrupts

- An operating system is **interrupt driven**
- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- Incoming interrupts are **disabled** while another interrupt is being processed to prevent a **lost interrupt**

Protection

Crossing the boundaries

- **System calls:** OS procedures that execute privileged instructions



- A **trap** is a software-generated interrupt caused either by an error or a user request

System calls for UNIX

System Call	Description
access()	This checks if a calling process has access to the required file
chdir()	The chdir command changes the current directory of the system
chmod()	The mode of a file can be changed using this command
chown()	This changes the ownership of a particular file
kill()	This system call sends kill signal to one or more processes
link()	A new file name is linked to an existing file using link system call.
open()	This opens a file for the reading or writing process
pause()	The pause call suspends a file until a particular signal occurs.
stime()	This system call sets the correct time.
times()	Gets the parent and child process times
alarm()	The alarm system call sets the alarm clock of a process
fork()	A new process is created using this command
chroot()	This changes the root directory of a file.
exit()	The exit system call is used to exit a process.


Memory protection

- **Architecture must provide support for the OS to:**
 - **Protect user programs from each other**
 - **Protect OS from the user programs**
- **Solution: Based register, Limit register**

Memory Structure


- **Main memory** - only large storage media that the CPU can access directly
- **Secondary storage** - extension of main memory that provides large nonvolatile storage capacity
- **Magnetic disks** - rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer

Storage-Device Hierarchy



Virtual memory

- Allows users to *run programs without loading it completely.*
- OS keeps track of the individual parts of the programs
- Translation lookaside buffers



ICS 212: Operating Systems


Lecture 3

Instructor:

Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

Abstract View of OS



Computer System

Registers

- Dedicated name for one word of memory in the CPU
- *General purpose* registers
- *Special purpose* registers

I/O Operation

- I/O devices and the CPU can execute concurrently
- Each *device controller* is in-charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- I/O techniques: *programmed I/O*, *Interrupt-driven I/O*, *direct memory access (DMA)*

Main OS Concepts

- ***Process*** Management
- ***Memory*** Management
- Information protection and security
- ***Scheduling*** and resource management
- System structure

Process Management

- **Process:** a *program in execution*
- **An instance of a program running on a computer**
- **The entity that can be assigned to and executed on a processor**
- **A unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources**

Memory Management

- **Process isolation**
- **Automatic allocation and management**
- **Support for modular programming**
- **Protection and access control**
- **Long-term storage**

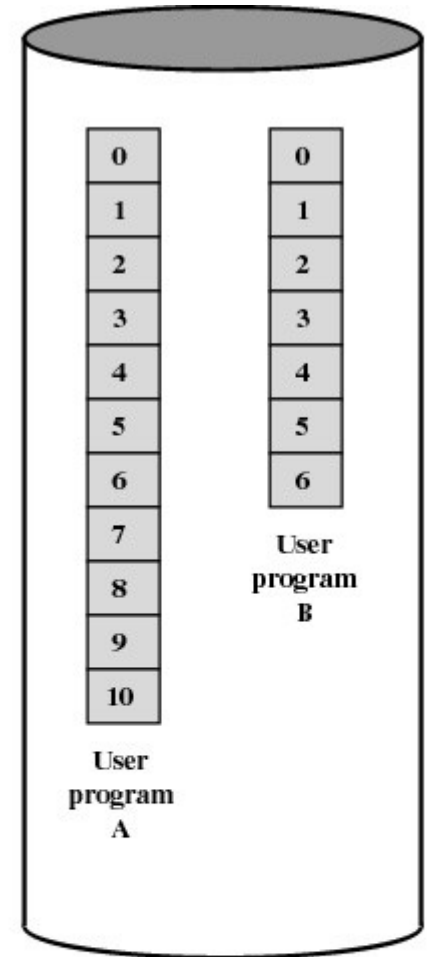
Paging

- Allows process to be comprised of a number of fixed-size blocks, called pages
- *Virtual address* is a page number and an offset within the page
- Each page may be located anywhere in main memory
- Real address or physical address in

A.1			
	A.0	A.2	
	A.5		
B.0	B.1	B.2	B.3
		A.7	
	A.9		
		A.8	
B.4	B.5	B.6	

Main Memory

Main memory consists of a number of fixed-length frames, equal to the size of a page. For a program to execute, some or all of its pages must be in main memory.



Disk

Secondary memory (disk) can hold many fixed-length pages. A user program consists of some number of pages. Pages for all programs plus the operating system are on disk, as are files.

Information Protection and Security

- ***Access control***
 - Regulate user access to the system
- ***Information flow control***
 - Regulate flow of data within the system and its delivery to users
- ***Certification***
 - Proving that access and flow control perform according to specifications

Scheduling and Resource Management

- ***Fairness***

- Give equal and fair access to all processes

- ***Differential responsiveness***

- Discriminate between different classes of jobs/users

- ***Efficiency***

- Maximize throughput, minimize response time, and accommodate as many uses as possible

ICS 212: Operating Systems

Lecture 4

Instructor:

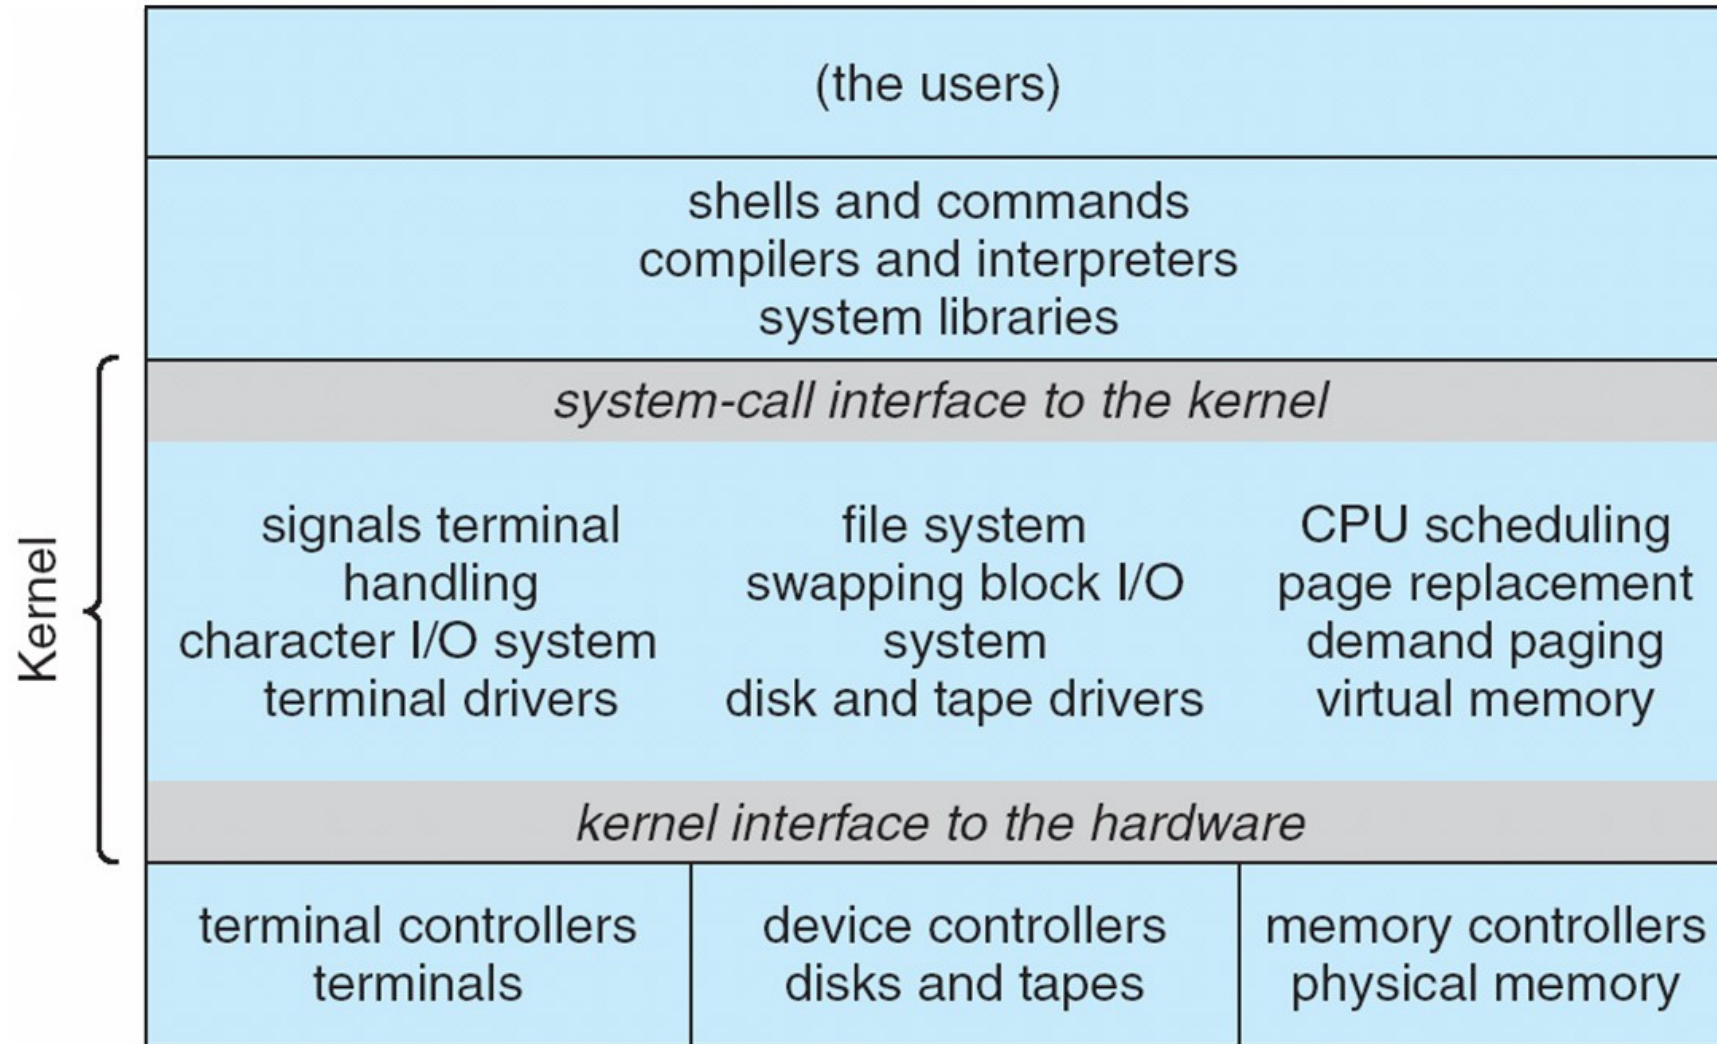
Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

OS Structure


- Kernel mode: user programs access via *system calls*
- Architectures:
 - *Monolithic kernel*
 - *Layered architecture*
 - *Microkernel*
 - *Modular approach*

Monolithic kernel



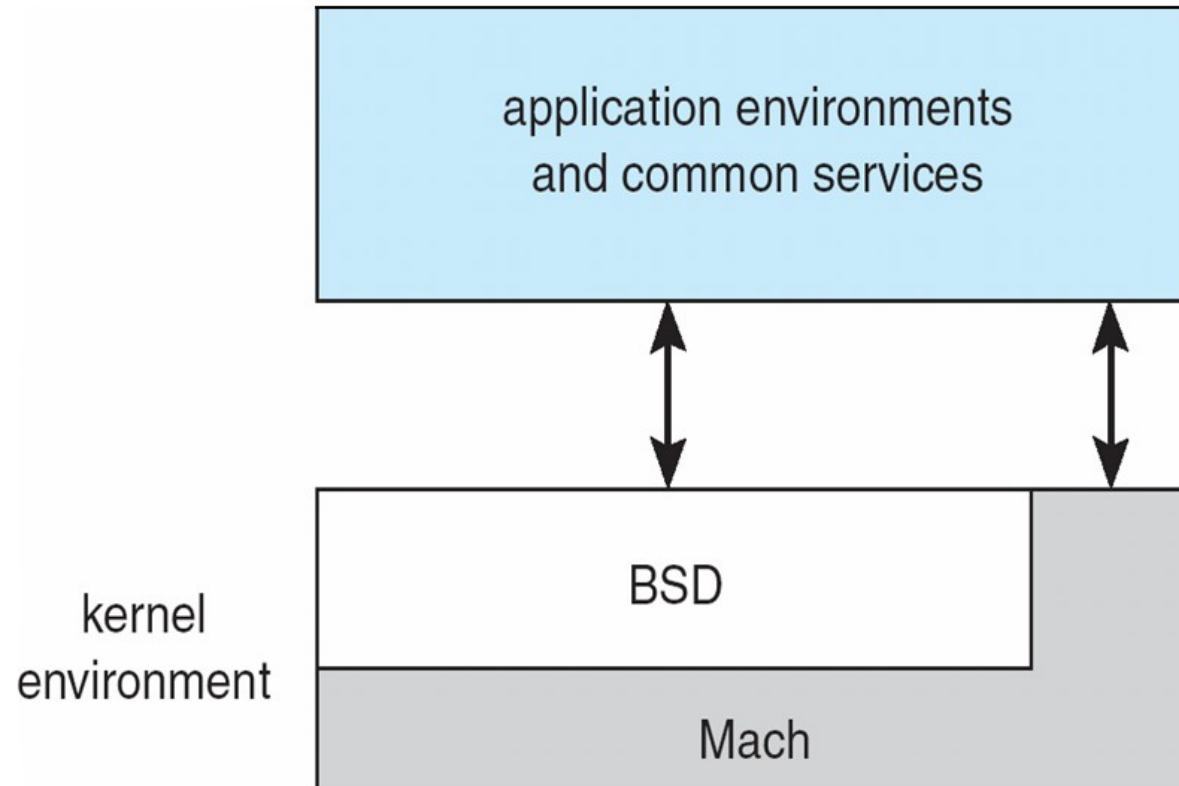
Layered architecture

- Layer uses Layer and provides services to Layer



Microkernel

- Moves as much from the kernel into “*user*” space



Benefits:

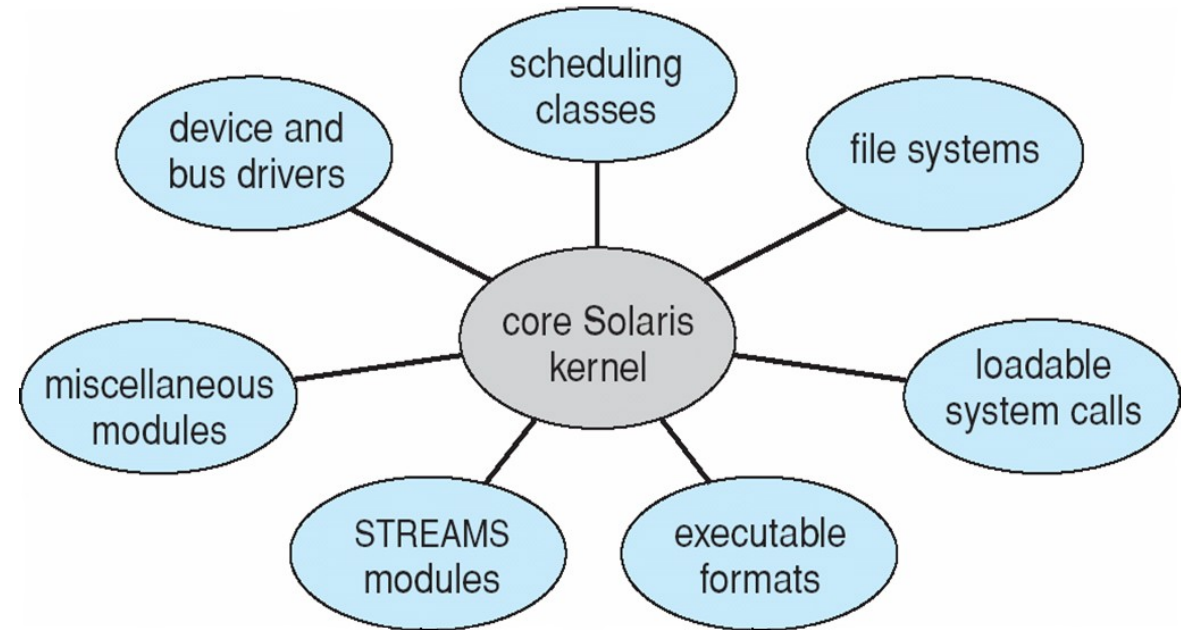
- Easier to extend a microkernel
- Easier to port the operating system to new architectures
- More reliable (less code is running in kernel mode)
- More secure

Disadvantage:

- Performance overhead of user space to kernel space communication

Modular approach


- **Most modern operating systems implement kernel modules**
 - **Uses object-oriented approach**
 - **Each core component is separate**
 - **Each talks to the others over known interfaces**
 - **Each is loadable as needed within the kernel**
- **Overall, similar to layers but with more flexible**



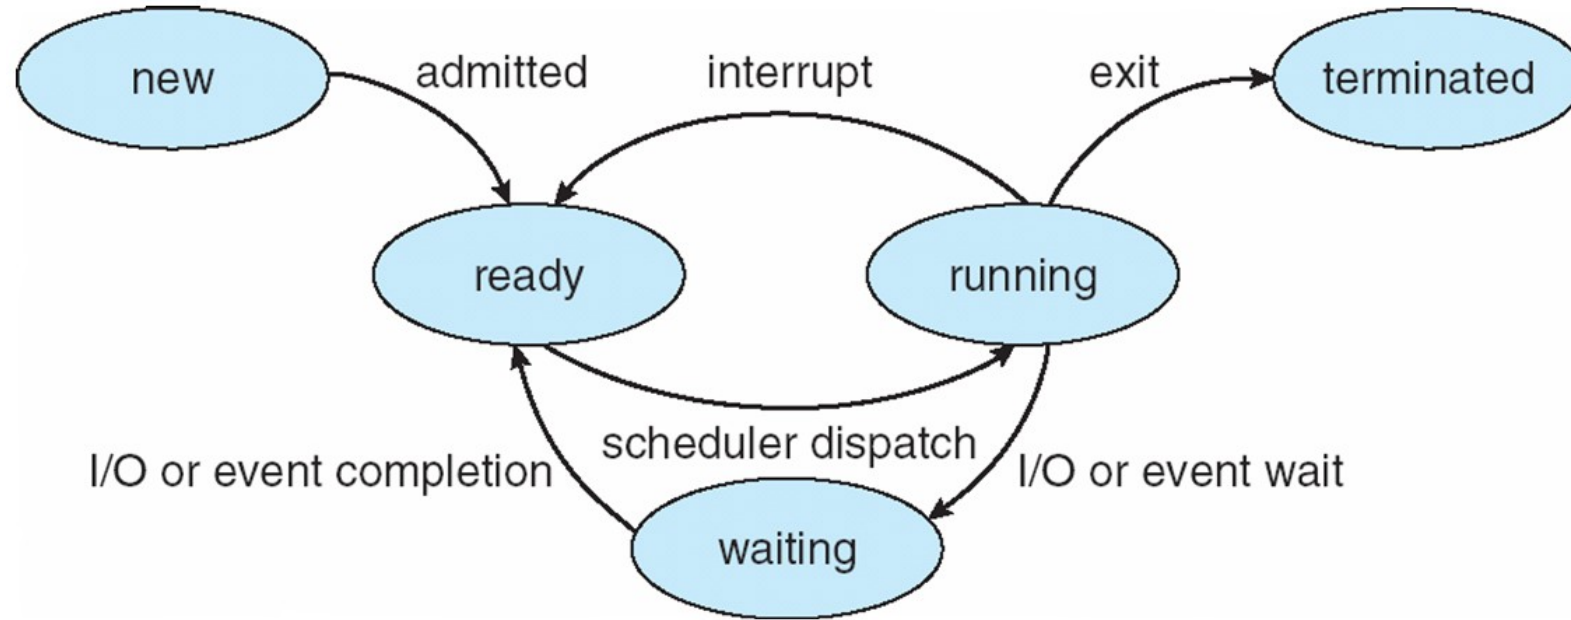
Processes

- An operating system executes a variety of programs:
 - Batch system - jobs
 - Time-shared systems - user programs or tasks
- Process - *a program in execution*
- Also called a *task*
- Can be traced
 - List the sequence of instructions that execute
- A process includes:
 - *Program counter*
 - *Stack pointer*
 - *Data section*
 - *etc.. (more on this soon)*

Process in Memory



Process state




Process Creation

- *Parent process create children processes*, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a *process identifier (pid)*
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation

- **Address space**
 - *Child duplicate of parent*
 - Child has a program loaded into it
- **UNIX examples**
 - **fork** system call creates new process
 - **exec** system call used after a fork to replace the process' memory space with a new program



Process Termination

- **Process executes last statement and asks the operating system to delete it (exit)**
 - **Output data from child to parent (via wait)**
 - **Process' resources are deallocated by operating system**
- **Parent may terminate execution of children processes (abort)**
 - **Child has exceeded allocated resources**
 - **Task assigned to child is no longer required**
 - ***If parent is exiting***
 - **Some operating system do not allow child to continue if its parent terminates**
 - **All children terminated - cascading termination**

ICS 212: Operating Systems

Lecture 5

Instructor:

Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

Last class

- **Process:** a *program in execution*
- Process States: *new, ready, executing, waiting, terminated*
- A process includes:
 - *program counter*
 - *stack*
 - *data section*

Process Control Block

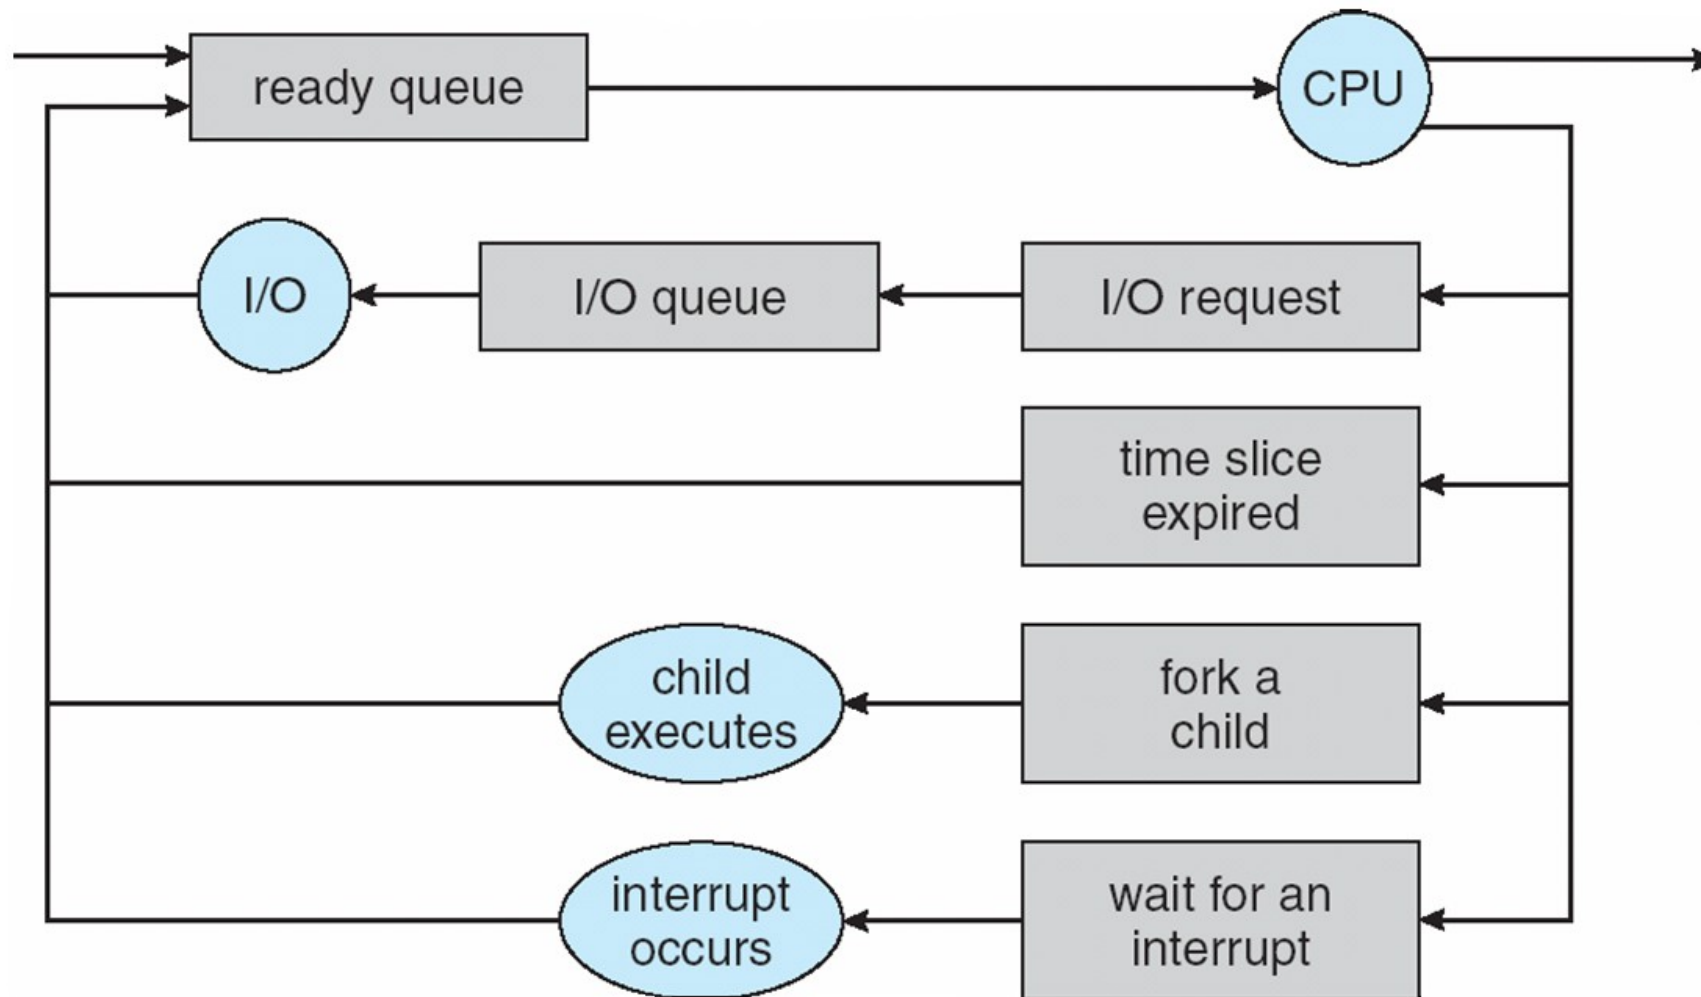
Information associated with each process

- *Process Identifier*
- *Process state*
- *Priority*
- *Program counter*
- *Memory pointers*
- *CPU scheduling information*
- *Accounting information*
- *I/O status information*

Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- *Processes migrate among the various queues*

Representing Process Scheduling

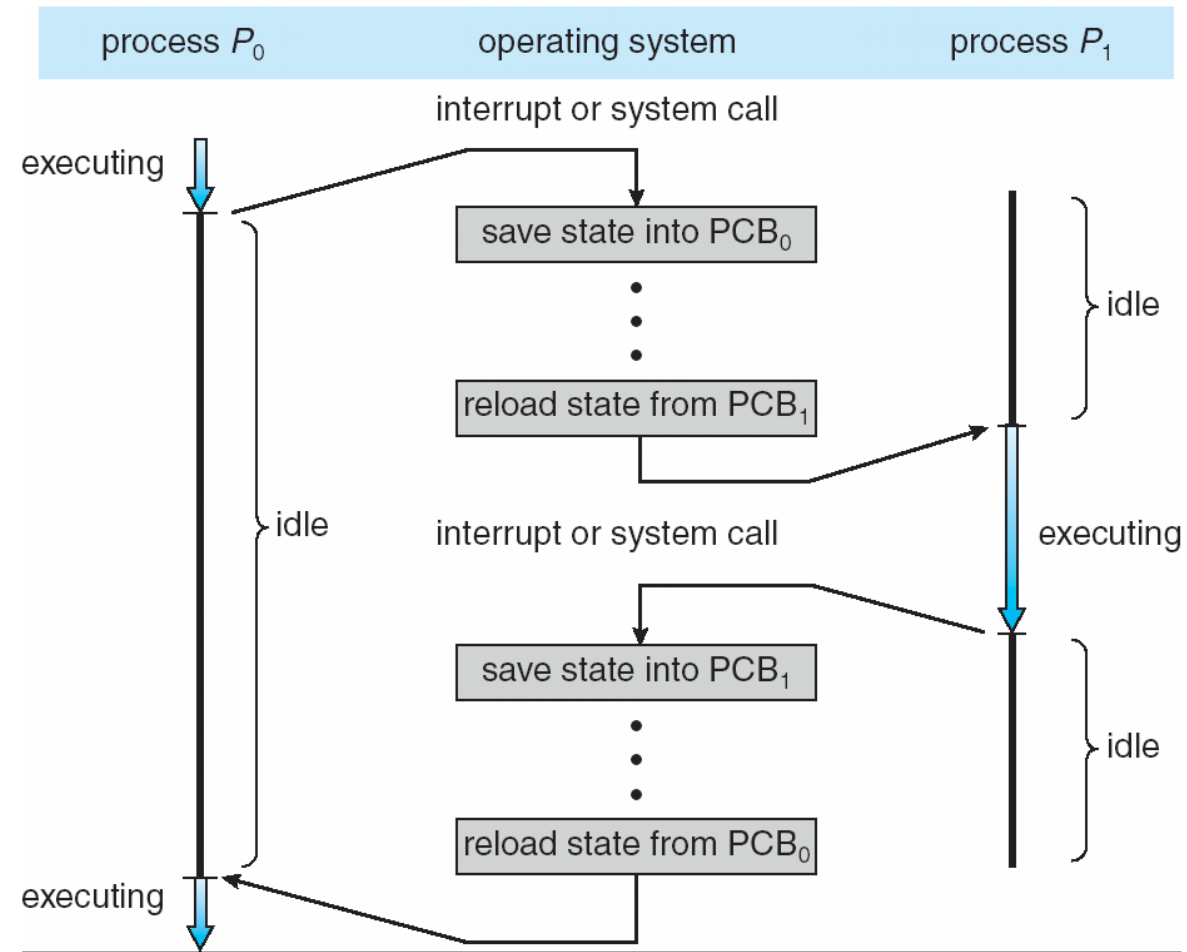


Process Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

Context Switching

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support



Independent and Cooperating Processes

- Processes within a system may be *independent* or *cooperating*
- Independent processes don't affect each other and don't share.
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperation:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
 - Message Passing
 - Shared memory

ICS 212: Operating Systems

Lecture 6

Instructor:

Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

Independent and Cooperating Processes

- Processes within a system may be *independent* or *cooperating*
- Independent processes don't affect each other and don't share.
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperation:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need *interprocess communication (IPC)*
 - Message Passing
 - Shared memory

ICS 212: Operating Systems

Lecture 7

Instructor:


Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

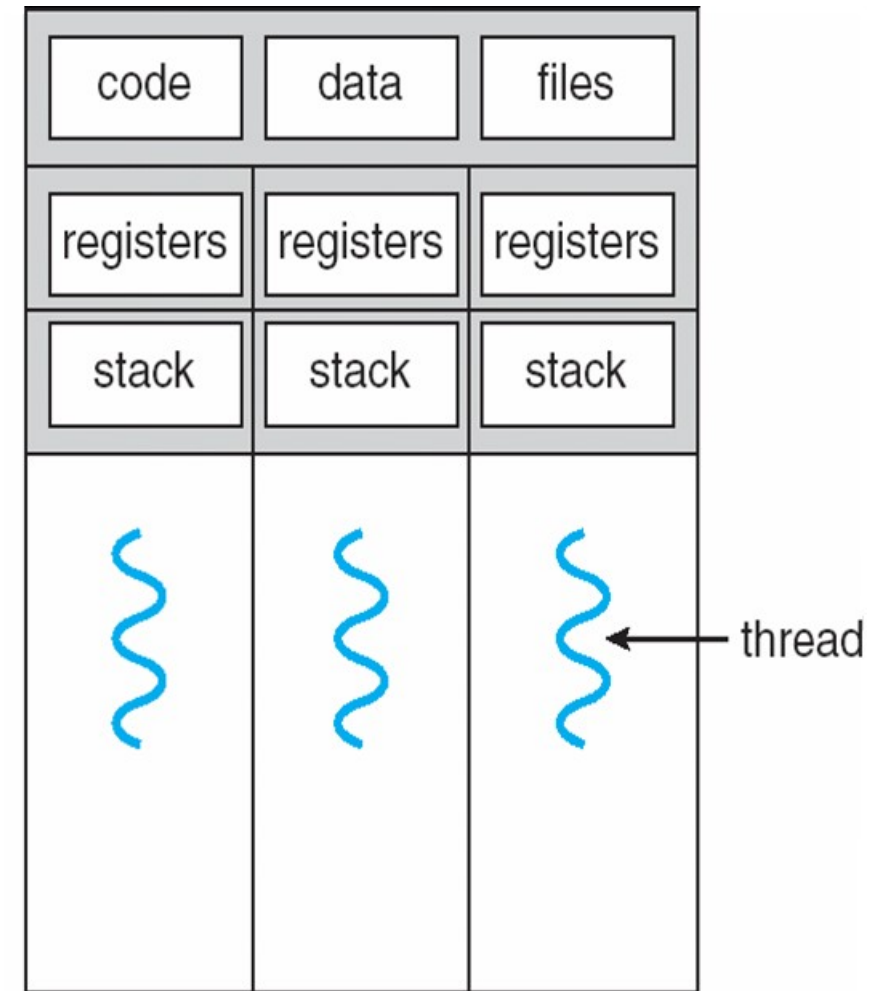
Threads

- A single sequential execution stream within a process.
- An ***independent unit of execution*** created ***within*** the context of a ***process***.
- ***Multithreading*** is a model of program execution that allows for ***multiple threads to be created within a process***, executing independently but concurrently sharing process resources.

Single threading v Multi threading



single-threaded process



multithreaded process

Why threading?

- Responsiveness
- Resource Sharing
- Economy
- Scalability

ICS 212: Operating Systems


Lecture 8

Instructor:

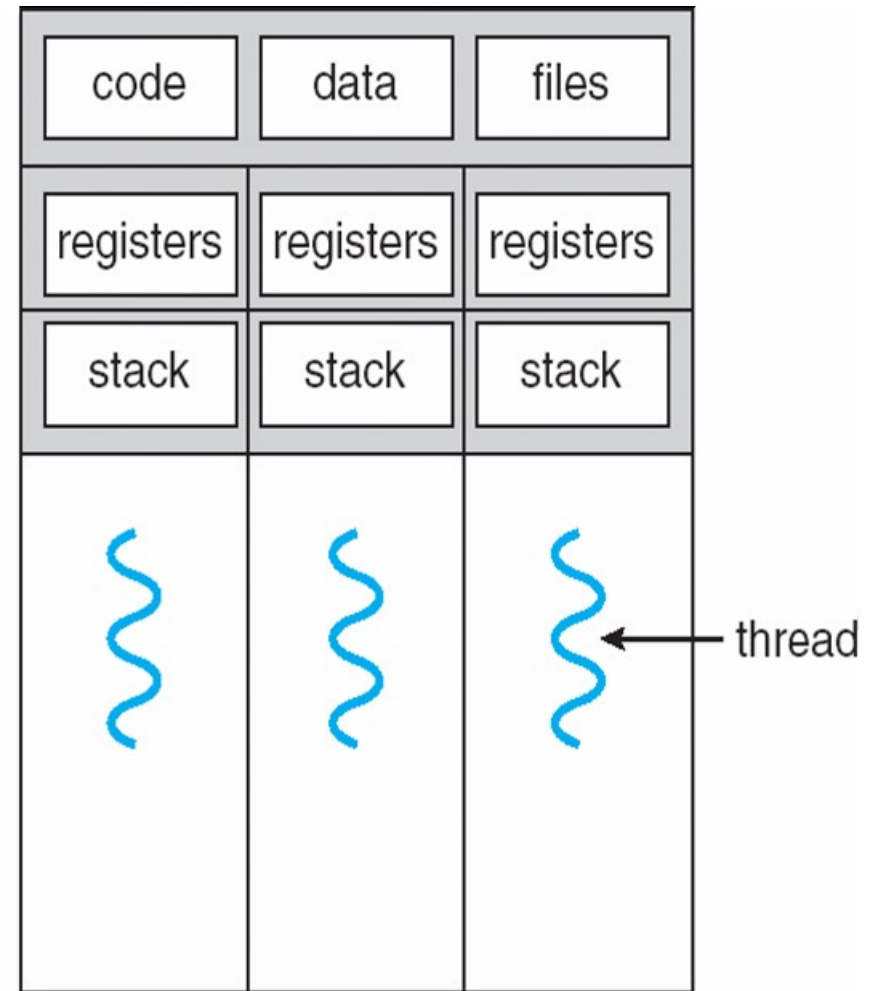
Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam


Processes and Threads




single-threaded process





multithreaded process





one process
one thread



one process
multiple threads



multiple processes
one thread per process



multiple processes
multiple threads per process

Types of threads

- User-level threads
 - *All thread management is done by the application*
 - The **kernel is not aware** of the existence of threads
- Advantages of user-level threads –
 - User-level threads are easier and faster to create than kernel-level threads.
 - They are more easily managed.
 - User-level threads can be run on any operating system.
 - There are no kernel mode privileges required for thread switching in user-level threads.

Types of threads

- **Kernel-level threads**
 - *Managed by the kernel*
 - Scheduling is done on a thread basis
- **Advantages of kernel-level threads:**
 - Multiple threads of the same process can be scheduled on different processors in kernel-level threads.
 - The kernel routines can also be multithreaded.
 - If a kernel-level thread is blocked, another thread of the same process can be scheduled by the kernel.

ICS 212: Operating Systems

Lecture 9

Instructor:


Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

Types of threads

- User-level threads
 - All thread management is done by the application
 - The kernel is not aware of the existence of threads
 - Implemented as a *thread library* which contains the code for thread creation, termination, scheduling and switching
- Kernel-level threads
 - Managed by the kernel
 - Scheduling is done on a thread basis


Many-to-one model




- Many user threads are mapped to one kernel thread
- The entire process will be **blocked** if a thread makes a blocking system call
- Multiple threads are unable to run in parallel

One-to-one model

- Each user thread is mapped to a kernel thread
- Multiple *threads can run in parallel* on multiprocessors
- Creating a user thread requires creating a corresponding kernel thread
- Implementational restriction




Many-to-many model



- Allows many user level threads to be mapped to some (equal or lesser) kernel threads
- Allows the *operating system to create a sufficient number of kernel threads*
- During a *blocking system call by one thread, kernel schedules another thread* for execution.

Two-level model



ICS 212: Operating Systems

Lecture 10

Instructor:

Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

Thread library

- Thread library provides programmers with APIs for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Common thread libraries

- POSIX threads:
 - A *POSIX* standard (IEEE 1003.1c) API
 - *Common in UNIX* based operating systems
 - Can be provided as *either a user level or a kernel level* library
- Win 32 threads
 - *For Windows* systems
 - *Only kernel* level library
- Java threads
 - Managed by the JVM
 - Typically implemented using the threads model provided by underlying OS as Java program

POSIX: Layers of Abstraction

- A *Pthread* is the POSIX specification for thread behaviour.
- A *kernel thread* is the entity actually created, scheduled and managed by the kernel as the context to execute a thread in a process
- “The *thread library may provide no more than a simple wrapper routine to a kernel system call to implement the POSIX specification*; or it may provide full support for thread creation, scheduling and management, independently of the kernel.”

Pthread objects

Type	Description
pthread_t	Thread identifier
Pthread_attr_t	Thread attribute
Pthread_cond_t	Condition variable
Pthread_condattr_t	Condition variable attribute
Pthread_key_t	Access key for thread-specific data
Pthread_mutex_t	mutex
Pthread_mutexattr_t	Mutex attribute
Pthread_once_t	One time initialization
Pthread_rwlock_t	Read-write lock
Pthread_rwlockattr_t	Read-write lock attribute

Thread management functions in POSIX

Function	Description
Pthread_create	Create a thread
Pthread_cancel	Terminate another thread
Pthread_detach	Set thread to release resources
Pthread_equal	Test two thread IDs for equality
Pthread_exit	Exit a thread w/o exiting process
Pthread_kill	Send a signal to a thread
Pthread_join	Wait for a thread

ICS 212: Operating Systems

Lecture 11

Instructor:

Dr. Nilotpal Chakraborty

Indian Institute of Information Technology Kottayam

